

```
1 import re
2 import numpy as np
3 import pandas as pd
4 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis as LDA
5 from sklearn.decomposition import PCA
6 from sklearn.decomposition import SparsePCA
7 from sklearn.tree import DecisionTreeClassifier
8 from sklearn.naive_bayes import BernoulliNB, GaussianNB
9 from imblearn.over_sampling import SMOTE
10 import pandas as pd
11 from sklearn.model_selection import train_test_split
12 import matplotlib.pyplot as plt
13 import random
14 from sklearn.preprocessing import StandardScaler
15 from sklearn.decomposition import TruncatedSVD
16 from sklearn.model_selection import GridSearchCV
17 from sklearn.metrics import accuracy_score, classification_report, confusion_matrix, f1_score, make_scorer
18 from sklearn.feature_selection import chi2
19 from sklearn.feature_selection import SelectKBest
20 from sklearn.impute import SimpleImputer
21 from sklearn.model_selection import ShuffleSplit, cross_val_score
```

```
1 bug2 = pd.read_csv("/content/drive/MyDrive/data_bug2.txt", sep=' ')
2 bug2.info()
3 print(bug2.shape)
4 bug2.head(7)
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 5 columns):
#   Column  Non-Null Count  Dtype
---  -
0    sx      100000 non-null    float64
1    sy      100000 non-null    float64
2    gx      100000 non-null    float64
3    gy      100000 non-null    float64
4    d        100000 non-null    float64
dtypes: float64(5)
memory usage: 3.8 MB
(100000, 5)
```

| | sx | sy | gx | gy | d |
|---|-----------|-----------|------------|----------|-------|
| 0 | -20.35270 | -13.16810 | -16.861800 | 13.85940 | 31.08 |

1 bug2.describe()

| | sx | sy | gx | gy | d |
|-------|---------------|---------------|---------------|---------------|---------------|
| count | 100000.000000 | 100000.000000 | 100000.000000 | 100000.000000 | 100000.000000 |
| mean | 2.558475 | -1.111631 | 2.461731 | -1.162655 | 62.395794 |
| std | 13.772926 | 9.125932 | 13.824945 | 9.114331 | 47.137511 |
| min | -21.999600 | -13.999900 | -21.999800 | -13.999900 | 0.000000 |
| 25% | -8.671688 | -9.578298 | -8.958005 | -9.627072 | 22.725000 |
| 50% | 4.375180 | -2.599365 | 4.251060 | -2.675130 | 55.020000 |
| 75% | 14.506300 | 7.541390 | 14.487900 | 7.334085 | 90.480000 |
| max | 21.999300 | 13.999900 | 21.999900 | 14.000000 | 258.660000 |

1 d_label = bug2['d']

1 bug2.drop(['d'],axis=1,inplace = True)

```
1 X_train, X_test, y_train, y_test = train_test_split(bug2, d_label, test_size=0.2, random_state=0)
```

```
1 from sklearn.preprocessing import StandardScaler
```

```
2 scaler = StandardScaler()
```

```
3 X_train1 = scaler.fit_transform(X_train)
```

```
4 print(scaler.mean_)
```

```
5 X_test1 = scaler.transform(X_test)
```

```
6
```

```
7 print(type(X_train1))
```

```
8 print(type(X_test1))
```

```
[ 2.5758674 -1.12390206  2.46324972 -1.16738687]
```

```
<class 'numpy.ndarray'>
```

```
<class 'numpy.ndarray'>
```

```
1 from sklearn.neighbors import KNeighborsRegressor
```

```
2 from sklearn.tree import DecisionTreeRegressor
```

```
3 from sklearn.ensemble import RandomForestRegressor
```

```
4 from sklearn.gaussian_process import GaussianProcessRegressor
```

```
5 from sklearn.neural_network import MLPRegressor
```

```
6 from sklearn.metrics import mean_squared_error
```

```
7 from sklearn.ensemble import AdaBoostRegressor
```

```
8
```

```
9 classifiers_list = {
```

```
10     "K-Neighbors Regressor": KNeighborsRegressor(),
```

```
11     "Decision Tree Regressor": DecisionTreeRegressor(),
```

```
12     "Random Forest Regressor": RandomForestRegressor(),
```

```
13     "Naive Bayes Regressor" : GaussianProcessRegressor(),
```

```
14     "Neural Network Regressor" : MLPRegressor(),
```

```
15     "Adaboost" : AdaBoostRegressor(base_estimator=RandomForestRegressor())
```

```
16 }
```

```
17
```

```
18 classifiers_count = len(classifiers_list.keys())
```

```
19 df_results = pd.DataFrame(data=np.zeros(shape=(classifiers_count,5)), columns = ['classifier', 'Recall', 'F1', 'Precision', 'Accuracy'])
```

```
20
```

```
21 for c_name, classifier in classifiers_list.items():
```

```
22     classifier.fit(X_train,y_train)
```

```
23     prediction = []
```

```
24     prediction = classifier.predict(X_test)
```

```
24 prediction = classifier.predict(x_test)
25 cv1 = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)
26 scores = cross_val_score(classifier, X_train, y_train, cv=cv1)
27 print ('Classifier+OneHotEncoder', c_name)
28 print ('Cross validation', scores)
29 print(prediction)
30 #print(classification_report(y_test,prediction))
31 print(mean_squared_error(y_test,prediction, squared=False))
```

```

1 classifiers_list = {
2     "K-Neighbors Regressor": KNeighborsRegressor(),
3     "Decision Tree Regressor": DecisionTreeRegressor(),
4     "Random Forest Regressor": RandomForestRegressor(),
5     #"Naive Bayes Regressor" : GaussianProcessRegressor(),
6     "Neural Network Regressor" : MLPRegressor()
7 }
8
9 classifiers_count = len(classifiers_list.keys())
10 df_results = pd.DataFrame(data=np.zeros(shape=(classifiers_count,5)), columns = ['classifier', 'Recall', 'F1', 'Precision', 'Accuracy'])
11
12 for c_name, classifier in classifiers_list.items():
13     classifier.fit(X_train1,y_train)
14     prediction = []
15     prediction = classifier.predict(X_test1)
16     cv1 = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)
17     scores = cross_val_score(classifier, X_train1, y_train, cv=cv1)
18     print ('Classifier+OneHotEncoder', c_name)
19     print ('Cross validation', scores)
20     print(prediction)
21     #print(classification_report(y_test,prediction))
22     print(mean_squared_error(y_test,prediction, squared=False))

```

```

Classifier+OneHotEncoder K-Neighbors Regressor
Cross validation [0.95790924 0.95540089 0.95564788 0.9566909  0.95506911]
[ 53.748  34.08 124.788 ... 45.108  25.392  54.24 ]
9.509439531623302

```

```

Classifier+OneHotEncoder Decision Tree Regressor
Cross validation [0.90688291 0.90322222 0.90317568 0.91455561 0.90675217]
[ 52.92  34.44 107.46 ... 45.36  26.46  52.86]
14.016947170479026

```

```

Classifier+OneHotEncoder Random Forest Regressor
Cross validation [0.95614924 0.95556907 0.95358417 0.95704907 0.95472031]
[ 51.354  34.7754 146.0046 ... 45.9414  25.239  52.6866]
9.693991643201576

```

```

/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma
ConvergenceWarning,
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma

```

```
ConvergenceWarning,  
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,  
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,  
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,  
Classifier+OneHotEncoder Neural Network Regressor  
Cross validation [0.81837745 0.79905087 0.82397237 0.82021138 0.8302946 ]  
[ 44.12417048 27.9305746 115.26066418 ... 24.98848521 18.61378319  
55.85931567]  
19.53060301989341  
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,
```

```
1 print(y_test)
```

```
3582      53.04  
60498     34.68  
53227    129.12  
21333     62.64  
3885      59.16  
  
...  
60116     61.98  
2415      20.94  
43763     45.84  
71345     24.90  
77687     53.16  
Name: d, Length: 20000, dtype: float64
```

```
1 import sklearn.metrics  
2 print(sklearn.metrics.SCORERS.keys())  
  
dict_keys(['explained_variance', 'r2', 'max_error', 'neg_median_absolute_error', 'neg_mean_absolute_error', 'neg_mean_absolute_percentage
```

```
1 from sklearn.model_selection import KFold  
2 from sklearn.model_selection import GridSearchCV
```

```

3 from sklearn.model_selection import train_test_split
4 from sklearn.model_selection import KFold
5
6 # Define our candidate hyperparameters
7 hp_candidates = [{'n_neighbors': [2,3,4,5,6,7,8,9,10,11,12,13,14,15], 'weights': ['uniform','distance'],'p':[1,2,5]}]
8
9 # Search for best hyperparameters
10 grid = GridSearchCV(estimator=KNeighborsRegressor(), param_grid=hp_candidates, cv=5,scoring='r2')
11
12 grid.fit(X_train1,y_train)

```

```

GridSearchCV(cv=5, estimator=KNeighborsRegressor(),
             param_grid=[{'n_neighbors': [2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12,
                                         13, 14, 15],
                           'p': [1, 2, 5], 'weights': ['uniform', 'distance']}],
             scoring='r2')

```

```

1 print("Tuned Hyperparameters :", grid.best_params_)
2 print("Accuracy :",grid.best_score_)

```

```

Tuned Hyperparameters : {'n_neighbors': 6, 'p': 2, 'weights': 'distance'}
Accuracy : 0.9582734325923965

```

```

1 hp_candidates = [{"max_depth" : [None,2,4,6,8,10,12], }]
2 grid = GridSearchCV(estimator=DecisionTreeRegressor(), param_grid=hp_candidates, cv=5)
3 grid.fit(X_train1,y_train)

```

```

GridSearchCV(cv=5, estimator=DecisionTreeRegressor(),
             param_grid=[{'max_depth': [None, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11,
                                         12]}])

```

```

1 print("Tuned Hyperparameters :", grid.best_params_)
2 print("Accuracy :",grid.best_score_)

```

```

Tuned Hyperparameters : {'max_depth': None}
Accuracy : 0.9057270160611951

```

```

1 hp_candidates = [{'n_estimators': [10,20,30,40],'max_features': ['auto', 'sqrt', 'log2'],'max_depth' : [None,2,3,4,5]}]
2 grid = GridSearchCV(estimator=RandomForestRegressor(), param_grid=hp_candidates, cv=5,scoring='neg mean absolute error')

```

```

3 grid.fit(X_train1,y_train)

    GridSearchCV(cv=5, estimator=RandomForestRegressor(),
        param_grid=[{'max_depth': [None, 2, 3, 4, 5],
            'max_features': ['auto', 'sqrt', 'log2'],
            'n_estimators': [10, 20, 30, 40]}],
        scoring='neg_mean_absolute_error')

1 print("Tuned Hyperparameters :", grid.best_params_)
2 print("Accuracy :",grid.best_score_)

    Tuned Hyperparameters : {'max_depth': None, 'max_features': 'auto', 'n_estimators': 40}
    Accuracy : -4.5210631125


1 classifiers_list = {
2     "K-Neighbors Regressor": KNeighborsRegressor(n_neighbors = 2, p = 2, weights = 'distance'),
3     "Decision Tree Regressor": DecisionTreeRegressor(max_depth= None),
4     "Random Forest Regressor": RandomForestRegressor(max_depth= None, max_features = 'auto', n_estimators = 40),
5     #"Naive Bayes Regressor" : GaussianProcessRegressor(),
6     "Neural Network Regressor" : MLPRegressor()
7 }
8
9 classifiers_count = len(classifiers_list.keys())
10 df_results = pd.DataFrame(data=np.zeros(shape=(classifiers_count,5)), columns = ['classifier', 'Recall', 'F1', 'Precision', 'Accuracy'])
11
12 for c_name, classifier in classifiers_list.items():
13     classifier.fit(X_train1,y_train)
14     prediction = []
15     prediction = classifier.predict(X_test1)
16     cv1 = ShuffleSplit(n_splits=5, test_size=0.2, random_state=0)
17     scores = cross_val_score(classifier, X_train1, y_train, cv=cv1)
18     print ('Classifier+OneHotEncoder', c_name)
19     print ('Cross validation', scores)
20     print(prediction)
21     #print(classification_report(y_test,prediction))
22     print(mean_squared_error(y_test,prediction, squared=False))

Classifier+OneHotEncoder K-Neighbors Regressor
Cross validation [0.95171266 0.9502014 0.95259018 0.95294602 0.94923148]

```



```
[ 52.30169881 32.99351873 145.48004059 ... 45.60925937 25.55738142
 52.83870022]
```

```
10.163535779197355
```

```
Classifier+OneHotEncoder Decision Tree Regressor
```

```
Cross validation [0.90723194 0.90185824 0.90154257 0.91268159 0.90455931]
```

```
[ 52.92 34.32 107.46 ... 45.36 26.46 52.74]
```

```
14.087015076303425
```

```
Classifier+OneHotEncoder Random Forest Regressor
```

```
Cross validation [0.9560381 0.95460111 0.95257575 0.95547944 0.95304388]
```

```
[ 51.828 34.5825 148.911 ... 45.585 25.4745 52.8315]
```

```
9.849071699470588
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,
```

```
Classifier+OneHotEncoder Neural Network Regressor
```

```
Cross validation [0.81797425 0.82557981 0.81994791 0.82640348 0.83188815]
```

```
[ 43.71138528 32.30164145 109.68240518 ... 33.30781509 27.30334201
```

```
64.86746883]
```

```
20.253207010597386
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:696: ConvergenceWarning: Stochastic Optimizer: Ma  
ConvergenceWarning,
```



```
1 from sklearn import neighbors
2 from sklearn.metrics import mean_squared_error
3 from math import sqrt
4 import matplotlib.pyplot as plt
5 %matplotlib inline
6 rmse_val = [] #to store rmse values for different k
7 for K in range(20):
8     K = K+1
9     model = neighbors.KNeighborsRegressor(n_neighbors = K)
10
11     model.fit(X_train, y_train) #fit the model
```

```
12     pred=model.predict(X_test) #make prediction on test set
13     error = sqrt(mean_squared_error(y_test,pred)) #calculate rmse
14     rmse_val.append(error) #store rmse values
15     print('RMSE value for k= ', K , 'is:', error)
```

```
RMSE value for k= 1 is: 11.888223827805398
RMSE value for k= 2 is: 10.265583827040722
RMSE value for k= 3 is: 9.842908972453214
RMSE value for k= 4 is: 9.722429523915306
RMSE value for k= 5 is: 9.655117643944067
RMSE value for k= 6 is: 9.639658604432006
RMSE value for k= 7 is: 9.677903698674436
RMSE value for k= 8 is: 9.722679412814916
RMSE value for k= 9 is: 9.768395860347
RMSE value for k= 10 is: 9.822937574096661
RMSE value for k= 11 is: 9.906196556234073
RMSE value for k= 12 is: 9.975580568380469
RMSE value for k= 13 is: 10.041009049219548
RMSE value for k= 14 is: 10.106092232527217
RMSE value for k= 15 is: 10.16928041611598
RMSE value for k= 16 is: 10.233886727844869
RMSE value for k= 17 is: 10.30946144081513
RMSE value for k= 18 is: 10.354768957103776
RMSE value for k= 19 is: 10.432988483980504
RMSE value for k= 20 is: 10.503287253376916
```

```
1 curve = pd.DataFrame(rmse_val) #elbow curve
2 curve.plot()
```

<matplotlib.axes._subplots.AxesSubplot at 0x7f87d92fb790>



```
1 hp_candidates = [{"max_depth" : [1,3,5,7,9,None],"min_samples_leaf":[2,3,4,5],"max_leaf_nodes": [None,10,20]}]
2 grid = GridSearchCV(estimator=DecisionTreeRegressor(), param_grid=hp_candidates, cv=5)
3 grid.fit(X_train1,y_train)
```

```
GridSearchCV(cv=5, estimator=DecisionTreeRegressor(),
             param_grid=[{'max_depth': [1, 3, 5, 7, 9, None],
                           'max_leaf_nodes': [None, 10, 20],
                           'min_samples_leaf': [2, 3, 4, 5]}])
```

```
1 print("Tuned Hyperparameters :", grid.best_params_)
2 print("Accuracy :",grid.best_score_)
```

```
Tuned Hyperparameters : {'max_depth': 9, 'max_leaf_nodes': None, 'min_samples_leaf': 3}
Accuracy : 0.8236125600489421
```

```
1 from sklearn.neighbors import KNeighborsRegressor
2 from sklearn.model_selection import KFold
3 from sklearn.model_selection import GridSearchCV
4 from sklearn.model_selection import train_test_split
5 from sklearn.model_selection import KFold
6
7 # Define our candidate hyperparameters
8 hp_candidates = [{'n_neighbors': [2,3,4,5,6,7,8,9,10,11,12,13,14,15], 'weights': ['uniform','distance'],'p':[1,2,5]}]
9
10 # Search for best hyperparameters
11 grid = GridSearchCV(estimator=KNeighborsRegressor(), param_grid=hp_candidates, cv=5)
12
13 grid.fit(X_train1,y_train)
```

```
1 print("Tuned Hyperparameters :", grid.best_params_)
2 print("Accuracy :",grid.best_score_)
```

```
1 rmse_val = [] #to store rmse values for different k
```

```
2 for K in range(20):
3     K = K+1
4     model = neighbors.KNeighborsRegressor(n_neighbors=K,p=2,weights='distance')
5
6     model.fit(X_train, y_train) #fit the model
7     pred=model.predict(X_test) #make prediction on test set
8     error = sqrt(mean_squared_error(y_test,pred)) #calculate rmse
9     rmse_val.append(error) #store rmse values
10    print('RMSE value for k= ' , K , 'is:', error)
```

```
1 curve = pd.DataFrame(rmse_val) #elbow curve
2 curve.plot()
```

```
1 from sklearn import tree
2 text_representation = tree.export_text(regr)
3 print(text_representation)
```

```
1 fig = plt.figure(figsize=(25,20))
2 _ = tree.plot_tree(regr, feature_names=['x','y','dx','dy'], filled=True)
```

[Colab paid products](#) - [Cancel contracts here](#)

▶ Executing (3m 24s) ... > ... > ... > _boo... > predi... > __call... > dispatch_one_b... > _dispat... > apply_as... > __init... > __call... > <listco... > __call... > _accumulate_predi... > predi... ... ✕